



# Chapter 3: Introduction to SQL

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Outline

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database



# History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory in the early 1970s
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
  - SQL-86
  - SQL-89
  - SQL-92
  - SQL:1999 (language name became Y2K compliant!)
  - SQL:2003, SQL:2006, SQL:2008, SQL:2011, SQL:2016



## History (Cont.)

- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
  - Not all examples here may work on your particular system



# Data Definition Language

- The SQL data-definition language (DDL) allows the specification of information about relations, including:
  - The schema for each relation.
  - The domain of values associated with each attribute.
  - Integrity constraints
  - And as we will see later, also other information such as
    - ▶ The set of indices to be maintained for each relations.
    - ▶ Security and authorization information for each relation.
    - ▶ The physical storage structure of each relation on disk.



# Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p*,*d*)**. Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point. (ex., **numeric**(3,1), allows 44.5 to be stored exactly, but not 444.5 or 0.32)
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.
- More are covered in Chapter 4.



# BASIC SCHEMA DEFINITION



# Create Table Construct

- An SQL relation is defined using the **create table** command:

```
create table  $r$  ( $A_1 D_1, A_2 D_2, \dots, A_n D_n,$   
                (integrity-constraint1),  
                ...,  
                (integrity-constraintk))
```

- $r$  is the name of the relation
- each  $A_i$  is an attribute name in the schema of relation  $r$
- $D_i$  is the data type of values in the domain of attribute  $A_i$





# Create Table Construct (Cont.)

## ■ Example:

```
create table instructor (  
    ID                char(5),  
    name              varchar(20),  
    dept_name         varchar(20),  
    salary           numeric(8,2)  
  
);
```



# Integrity Constraints in Create Table

- **primary key** ( $A_1, \dots, A_n$ )
- **foreign key** ( $A_m, \dots, A_n$ ) **references**  $r$
- **not null**

*Example:*

```
create table instructor (  
    ID          char(5),  
    name        varchar(20) not null,  
    dept_name varchar(20),  
    salary      numeric(8,2),  
    primary key (ID),  
    foreign key (dept_name) references department);
```

**primary key** declaration on an attribute automatically ensures **not null**



## And a Few More Relation Definitions

```
create table student (  
    ID                varchar(5),  
    name             varchar(20) not null,  
    dept_name        varchar(20),  
    tot_cred         numeric(3,0),  
    primary key (ID),  
    foreign key (dept_name) references  
    department);
```



## And a Few More Relation Definitions (Cont.)

```
create table takes (  
    ID          varchar(5),  
    course_id  varchar(8),  
    sec_id     varchar(8),  
    semester   varchar(6),  
    year       numeric(4,0),  
    grade      varchar(2),  
    primary key (ID, course_id, sec_id, semester,  
year) ,  
    foreign key (ID) references student,  
    foreign key (course_id, sec_id, semester,  
year) references section);
```



## And more still

```
create table course (  
    course_id      varchar(8),  
    title          varchar(50),  
    dept_name     varchar(20),  
    credits        numeric(2,0),  
    primary key (course_id),  
    foreign key (dept_name) references  
    department);
```



# Updates to tables

## ■ Insert

- **insert into** *instructor* **values** ('10211', 'Smith', 'Biology', 66000);

## ■ Delete

- Remove all tuples from the *student* relation
  - ▶ **delete from** *student*

## ■ Drop Table

- **drop table** *r*

## ■ What is the difference between delete and drop?



# Updates to tables (Cont.)

## ■ Alter

### ● **alter table $r$ add $A$ $D$**

- ▶ where  $A$  is the name of the attribute to be added to relation  $r$  and  $D$  is the domain of  $A$ .
- ▶ All exiting tuples in the relation are assigned *null* as the value for the new attribute.

### ● **alter table $r$ drop $A$**

- ▶ where  $A$  is the name of an attribute of relation  $r$
- ▶ Dropping of attributes not supported by many databases. Only allow to drop the whole table.
- ▶ Can also drop constraint



# Basic Query Structure

- A typical SQL query has the form:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

- $A_i$  represents an attribute
  - $R_i$  represents a relation
  - $P$  is a predicate.
- The result of an SQL query is a relation.
    - But, may have duplicates, not a relation according to formal mathematical definition





# The select Clause

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:  
**select** *name*  
**from** *instructor*



# The select Clause (Cont.)

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after select.
- Find the department names of all instructors, and remove duplicates

```
select distinct dept_name  
from instructor
```

- The keyword **all** specifies that duplicates should not be removed.

```
select all dept_name  
from instructor
```



# The select Clause (Cont.)

- An asterisk in the select clause denotes “all attributes”

**select \***  
**from** *instructor*

| <i>name</i> |
|-------------|
| Srinivasan  |
| Wu          |
| Mozart      |
| Einstein    |
| El Said     |
| Gold        |
| Katz        |
| Califieri   |
| Singh       |
| Crick       |
| Brandt      |
| Kim         |

**Figure 3.2** Result of “select *name* from *instructor*”.

| <i>dept_name</i> |
|------------------|
| Comp. Sci.       |
| Finance          |
| Music            |
| Physics          |
| History          |
| Physics          |
| Comp. Sci.       |
| History          |
| Finance          |
| Biology          |
| Comp. Sci.       |
| Elec. Eng.       |

**Figure 3.3** Result of “select *dept\_name* from *instructor*”.



## The select Clause (Cont.)

- The **select** clause can contain arithmetic expressions involving the operation, +, −, \*, and /, and operating on constants or attributes of tuples.

- The query:

```
select ID, name, salary/12  
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

- Can rename “*salary/12*” using the **as** clause:

```
select ID, name, salary/12 as monthly_salary
```



# The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

```
select name
from instructor
where dept_name = 'Comp. Sci.'
```



## The where Clause (Cont.)

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
  - To find all instructors in Comp. Sci. dept with salary > 80000

```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary >
80000
```

- Comparisons can be applied to results of arithmetic expressions.



# The from Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.

- Find the Cartesian product *instructor X teaches*

**select \***  
**from** *instructor, teaches*

- generates every possible instructor – teaches pair, with all attributes from both relations.
  - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra).



# Cartesian Product

*instructor*

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 12121     | Wu          | Finance          | 90000         |
| 15151     | Mozart      | Music            | 40000         |
| 22222     | Einstein    | Physics          | 95000         |
| 32343     | El Said     | History          | 60000         |

*teaches*

| <i>ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> |
|-----------|------------------|---------------|-----------------|-------------|
| 10101     | CS-101           | 1             | Fall            | 2009        |
| 10101     | CS-315           | 1             | Spring          | 2010        |
| 10101     | CS-347           | 1             | Fall            | 2009        |
| 12121     | FIN-201          | 1             | Spring          | 2010        |
| 15151     | MU-199           | 1             | Spring          | 2010        |
| 22222     | PHY-101          | 1             | Fall            | 2009        |

| <i>Inst.ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> | <i>teaches.ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> |
|----------------|-------------|------------------|---------------|-------------------|------------------|---------------|-----------------|-------------|
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 10101             | CS-101           | 1             | Fall            | 2009        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 10101             | CS-315           | 1             | Spring          | 2010        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 10101             | CS-347           | 1             | Fall            | 2009        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 12121             | FIN-201          | 1             | Spring          | 2010        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 15151             | MU-199           | 1             | Spring          | 2010        |
| 10101          | Srinivasan  | Comp. Sci.       | 65000         | 22222             | PHY-101          | 1             | Fall            | 2009        |
| ...            | ...         | ...              | ...           | ...               | ...              | ...           | ...             | ...         |
| ...            | ...         | ...              | ...           | ...               | ...              | ...           | ...             | ...         |
| 12121          | Wu          | Finance          | 90000         | 10101             | CS-101           | 1             | Fall            | 2009        |
| 12121          | Wu          | Finance          | 90000         | 10101             | CS-315           | 1             | Spring          | 2010        |
| 12121          | Wu          | Pinance          | 90000         | 10101             | CS-347           | 1             | Fall            | 2009        |
| 12121          | Wu          | Pinance          | 90000         | 12121             | FIN-201          | 1             | Spring          | 2010        |
| 12121          | Wu          | Finance          | 90000         | 15151             | MU-199           | 1             | Spring          | 2010        |
| 12121          | Wu          | Pinance          | 90000         | 22222             | PHY-101          | 1             | Fall            | 2009        |
| ...            | ...         | ...              | ...           | ...               | ...              | ...           | ...             | ...         |
| ...            | ...         | ...              | ...           | ...               | ...              | ...           | ...             | ...         |





# Examples

- Find the names of all instructors who have taught some course and the course\_id

```
select name, course_id  
from instructor, teaches  
where instructor.ID = teaches.ID
```

- Find the names of all instructors in the Art department who have taught some course and the course\_id

```
select name, course_id  
from instructor , teaches  
where instructor.ID = teaches.ID and instructor.  
dept_name = 'Art'
```



# Understanding SQL Query

- In general, the meaning of an SQL query can be understood as follows:
  - Generate a Cartesian product of the relations listed in the **from** clause
  - Apply the predicates specified in the **where** clause on the result of Step 1
  - For each tuple in the result of Step 2, output the attributes (or results of expressions) specified in the **select** clause

Note: The above sequence of steps helps make clear what the result of an SQL query should be, **not** how it should be executed. A real implementation of SQL optimize evaluation by generating (as far as possible) only elements of the Cartesian product that satisfy the **where** clause predicates.



# The Rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

*old-name as new-name*

- Example:

```
select name as instructor_name, course_id  
from instructor, teaches  
where instructor.ID = teaches.ID;
```

```
select T.name, S.course_id  
from instructor as T, teaches as S  
where T.ID = S.ID;
```



# Join

- Matches attributes with the same value
- Consider the query “For all instructors in the university who have taught some course, find their names and the course ID of all courses they taught”, we can write as

**select** *name, course id*

**from** *instructor, teaches*

**where** *instructor.ID = teaches.ID;*



## Join (Cont.)

- It is equivalent to  
**select** *name, course id*  
**from** *instructor* **join** *teaches*  
**on** *instructor.ID = teaches.ID;*



# String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring.
  - underscore ( \_ ). The \_ character matches any character.



## String Operations (Cont.)

- Find the names of all instructors whose name includes the substring “dar”.

```
select name  
from instructor  
where name like '%dar%'
```

- Match the string “100%”

```
like '100 \%' escape '\'
```

in that above we use backslash (\) as the escape character.



# String Operations (Cont.)

- Patterns are case sensitive.
- Pattern matching examples:
  - 'Intro%' matches any string beginning with "Intro".
  - '%Comp%' matches any string containing "Comp" as a substring.
  - '\_\_\_' matches any string of exactly three characters.
  - '\_\_\_ %' matches any string of at least three characters.





# String Operations (Cont.)

- SQL supports a variety of string operations such as
  - Concatenation (CONCAT)
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.

<https://docs.microsoft.com/en-us/sql/t-sql/functions/string-functions-transact-sql?view=sql-server-2017>



# Ordering the Display of Tuples

- List in alphabetic order the names of all instructors

```
select distinct name  
from instructor  
order by name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

- Example: **order by** *name* **desc**

- Can sort on multiple attributes

- Example: **order by** *dept\_name*, *name*

- **order by** *salary* **desc**, *name* **asc**;



# Where Clause Predicates

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq$  \$90,000 and  $\leq$  \$100,000)

```
select name  
from instructor  
where salary between 90000 and 100000
```



## ■ Consider the following problems

- Find courses (just course id information) that ran in Fall 2009 or in Spring 2010
- Find courses that ran in both Fall 2009 and Spring 2010
- Find courses that ran in Fall 2009 but not in Spring 2010



# Set Operations

- Find courses that ran in Fall 2009 or in Spring 2010

**(select *course\_id* from *section* where *semester* = 'Fall' and *year* = 2009)**  
**union**

**(select *course\_id* from *section* where *semester* = 'Spring' and *year* = 2010)**

- Find courses that ran in both Fall 2009 and Spring 2010

**(select *course\_id* from *section* where *semester* = 'Fall' and *year* = 2009)**  
**intersect**

**(select *course\_id* from *section* where *semester* = 'Spring' and *year* = 2010)**

- Find courses that ran in Fall 2009 but not in Spring 2010

**(select *course\_id* from *section* where *semester* = 'Fall' and *year* = 2009)**  
**except**

**(select *course\_id* from *section* where *semester* = 'Spring' and *year* = 2010)**



## Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary.

**(select distinct *salary***  
**from *instructor*)**

**except**

**(select max(salary)**  
**from *instructor*)**



# Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```



# Null Values and Three Valued Logic

- Three values – *true*, *false*, *unknown*
- Any comparison with *null* returns *unknown*
  - Example:  $5 < null$  or  $null <> null$  or  $null = null$
- Three-valued logic using the value *unknown*:
  - OR:  $(unknown \text{ or } true) = true$ ,  
 $(unknown \text{ or } false) = unknown$   
 $(unknown \text{ or } unknown) = unknown$
  - AND:  $(true \text{ and } unknown) = unknown$ ,  
 $(false \text{ and } unknown) = false$ ,  
 $(unknown \text{ and } unknown) = unknown$
  - NOT:  $(\text{not } unknown) = unknown$
  - “*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*





# Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values



# Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department

```
select avg (salary)  
from instructor  
where dept_name= 'Comp. Sci.';
```

- Find the total number of instructors who teach a course in the Spring 2010 semester

```
select count (distinct ID)  
from teaches  
where semester = 'Spring' and year = 2010;
```

- Find the number of tuples in the *course* relation

```
select count (*)  
from course;
```



# Aggregate Functions – Group By

- Find the average salary of instructors in each department
  - **select** *dept\_name*, **avg** (*salary*) **as** *avg\_salary*  
**from** *instructor*

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 76766     | Crick       | Biology          | 72000         |
| 45565     | Katz        | Comp. Sci.       | 75000         |
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 83821     | Brandt      | Comp. Sci.       | 92000         |
| 98345     | Kim         | Elec. Eng.       | 80000         |
| 12121     | Wu          | Finance          | 90000         |
| 76543     | Singh       | Finance          | 80000         |
| 32343     | El Said     | History          | 60000         |
| 58583     | Califieri   | History          | 62000         |
| 15151     | Mozart      | Music            | 40000         |
| 33456     | Gold        | Physics          | 87000         |
| 22222     | Einstein    | Physics          | 95000         |

| <i>dept_name</i> | <i>avg_salary</i> |
|------------------|-------------------|
| Biology          | 72000             |
| Comp. Sci.       | 77333             |
| Elec. Eng.       | 80000             |
| Finance          | 85000             |
| History          | 61000             |
| Music            | 40000             |
| Physics          | 91000             |



## Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list
  - */\* erroneous query \*/*  
**select** *dept\_name*, *ID*, **avg** (*salary*)  
**from** *instructor*  
**group by** *dept\_name*;



# Aggregate Functions – Having Clause

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups



# Null Values and Aggregates

## ■ Total all salaries

```
select sum (salary )  
from instructor
```

- Above statement ignores null amounts
  - Result is *null* if there is no non-null amount
- ## ■ All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- ## ■ What if collection has only null values?
- count returns 0
  - all other aggregates return null



# Exercise

- Using the files provided to generate the University schema
- Write the following queries
  - Find the titles of courses in the Comp. Sci. department that have 3 credits.
  - Find the highest salary of any instructor.
  - Find all instructors earning the highest salary
  - Find the enrollment of each section that was offered in Spring 2009



# Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.





# Subqueries in the Where Clause



# Subqueries in the Where Clause

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality



# Set Membership

- Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year = 2009 and
       course_id in (select course_id
                       from section
                       where semester = 'Spring'
                       and year = 2010);
```



## Set Membership (Cont.)

- Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id  
from section  
where semester = 'Fall' and year = 2009 and  
       course_id not in (select course_id  
                           from section  
                           where semester = 'Spring' and  
                           year = 2010);
```



## Set Comparison – “some” Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Physics department.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept_name = Physics';
```

Same query using > **some** clause

```
select name  
from instructor  
where salary > some (select salary  
                        from instructor  
                        where dept_name = 'Physics');
```



# Definition of “some” Clause

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$  (read:  $5 < \text{some tuple in the relation}$ )

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However,  $(\neq \text{some}) \not\equiv \text{not in}$



## Set Comparison – “all” Clause

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Computer Science department.

```
select name  
from instructor  
where salary > all (select salary  
                        from instructor  
                        where dept_name = 'Comp. Sci.');
```



# Definition of “all” Clause

$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$

$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$

$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$

$(\neq \text{all}) \equiv \text{not in}$

However,  $(= \text{all}) \not\equiv \text{in}$





# Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$



# Use of “exists” Clause

- Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
      exists (select *
              from section as T
              where semester = 'Spring' and year = 2010
              and S.course_id = T.course_id);
```

- **Correlation name** – variable *S* in the outer query
- **Correlated subquery** – the inner query



# Use of “exists” Clause

- Find all courses taught in the Fall 2009 semester but not in the Spring 2010 semester.



# Subqueries in the From Clause



# Subqueries in the From Clause

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary  
from (select dept_name, avg (salary) as avg_salary  
      from instructor  
      group by dept_name) as dept_average  
where avg_salary > 42000;
```

- Note that we do not need to use the **having** clause



## With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the budget over the average  
**with** *avg\_budget(value)* **as**  
    (**select avg**(*budget*)  
    **from** *department*)  
**select** *department.dept\_name*  
**from** *department, avg\_budget*  
**where** *department.budget* > *avg\_budget.value*;



# Subqueries in the Select Clause



# Subquery in SELECT

- List all departments along with the number of instructors in each department

```
select dept_name,  
      (select count(*)  
      from instructor  
      where department.dept_name = instructor.dept_name)  
      as num_instructors  
  
from department;
```





# Modification of the Database

- Deletion of tuples from a given relation.
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation



# Deletion

- Delete all instructors

**delete from** *instructor*

- Delete all instructors from the Finance department

**delete from** *instructor*  
**where** *dept\_name* = 'Finance';

- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

**delete from** *instructor*  
**where** *dept name* in (**select** *dept name*  
**from** *department*  
**where** *building* = 'Watson');



## Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors  
**delete from** *instructor*  
**where** *salary* < (**select avg** (*salary*)  
                                  **from** *instructor*);
- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL (SQL already solves this problem by the following processing):
  1. Compute **avg** (*salary*) and find all tuples to delete
  2. Delete all tuples found above (without recomputing **avg** or retesting the tuples)



# Insertion

- Add a new tuple to *course*

**insert into** *course* **values** ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- or equivalently

**insert into** *course* (*course\_id*, *title*, *dept\_name*, *credits*) **values** ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- Add a new tuple to *student* with *tot\_creds* set to null

**insert into** *student*  
**values** ('3003', 'Green', 'Finance', *null*);



# Updates

- Increase salaries of instructors whose salary is over \$90,000 by 3%, and all others by a 5%
  - Write two **update** statements:  

```
update instructor
  set salary = salary * 1.03
  where salary > 90000;
update instructor
  set salary = salary * 1.05
  where salary <= 90000;
```
  - The order is important (what will happen if we change the order of the two update blocks?)
  - Can be done better using the **case** statement (next slide)



# Case Statement for Conditional Updates

- Same query as before but with case statement

**update** *instructor*

**set** *salary* = **case**

**when** *salary* <= 90000

**then** *salary* \* 1.05

**else** *salary* \* 1.03

**end**



## Exercise

- Find the names of all students who have taken at least one Comp. Sci. course; make sure there are no duplicate names in the result
- Create a new course “CS-001”, titled “Weekly Seminar”, with 0 credits.
- Increase the salary of each instructor in the Comp. Sci. department by 10%.
- Insert every student whose tot cred attribute is greater than 100 as an instructor in the same department, with a salary of 30000.



## Exercise

- Write SQL DDL corresponding to the following schema. Make any reasonable assumptions about data types, and be sure to declare primary and foreign keys.

*person* (*driver\_id*, *name*, *address*)

*car* (*license*, *model*, *year*)

*accident* (*report\_number*, *date*, *location*)

*owns* (*driver\_id*, *license*)

*participated* (*report\_number*, *license*, *driver\_id*, *damage\_amount*)





*person* (*driver\_id*, *name*, *address*)

*car* (*license*, *model*, *year*)

*accident* (*report\_number*, *date*, *location*)

*owns* (*driver\_id*, *license*)

*participated* (*report\_number*, *license*, *driver\_id*, *damage\_amount*)

**create table** *person*

(*driver\_id* **varchar**(50),

*name* **varchar**(50),

*address* **varchar**(50),

**primary key** (*driver id*))



*person* (*driver\_id*, *name*, *address*)

*car* (*license*, *model*, *year*)

*accident* (*report\_number*, *date*, *location*)

*owns* (*driver\_id*, *license*)

*participated* (*report\_number*, *license*, *driver\_id*, *damage\_amount*)

**create table** *car*

(*license* **varchar**(50),

*model* **varchar**(50),

*year* **integer**,

**primary key** (*license*))



*person* (*driver\_id*, *name*, *address*)

*car* (*license*, *model*, *year*)

*accident* (*report\_number*, *date*, *location*)

*owns* (*driver\_id*, *license*)

*participated* (*report\_number*, *license*, *driver\_id*, *damage\_amount*)

**create table** *accident*

(*report\_number* **integer**,

*date* **date**,

*location* **varchar**(50),

**primary key** (*report\_number*))



*person* (*driver\_id*, *name*, *address*)

*car* (*license*, *model*, *year*)

*accident* (*report\_number*, *date*, *location*)

*owns* (*driver\_id*, *license*)

*participated* (*report\_number*, *license*, *driver\_id*, *damage\_amount*)

**create table** *owns*

(*driver\_id* **varchar**(50),

*license* **varchar**(50),

**primary key** (*driver\_id*, *license*)

**foreign key** (*driver\_id*) **references** *person*

**foreign key** (*license*) **references** *car*)



**create table** *participated*

*(report\_number* **integer**,

*license* **varchar**(50),

*driver\_id* **varchar**(50),

*damage\_amount* **integer**,

**primary key** (*report\_number*, *license*)

**foriegn key** (*license*) **references** *car*

**foriegn key** (*report number*) **references** *accident*))