

MongoDB-Advanced Topics

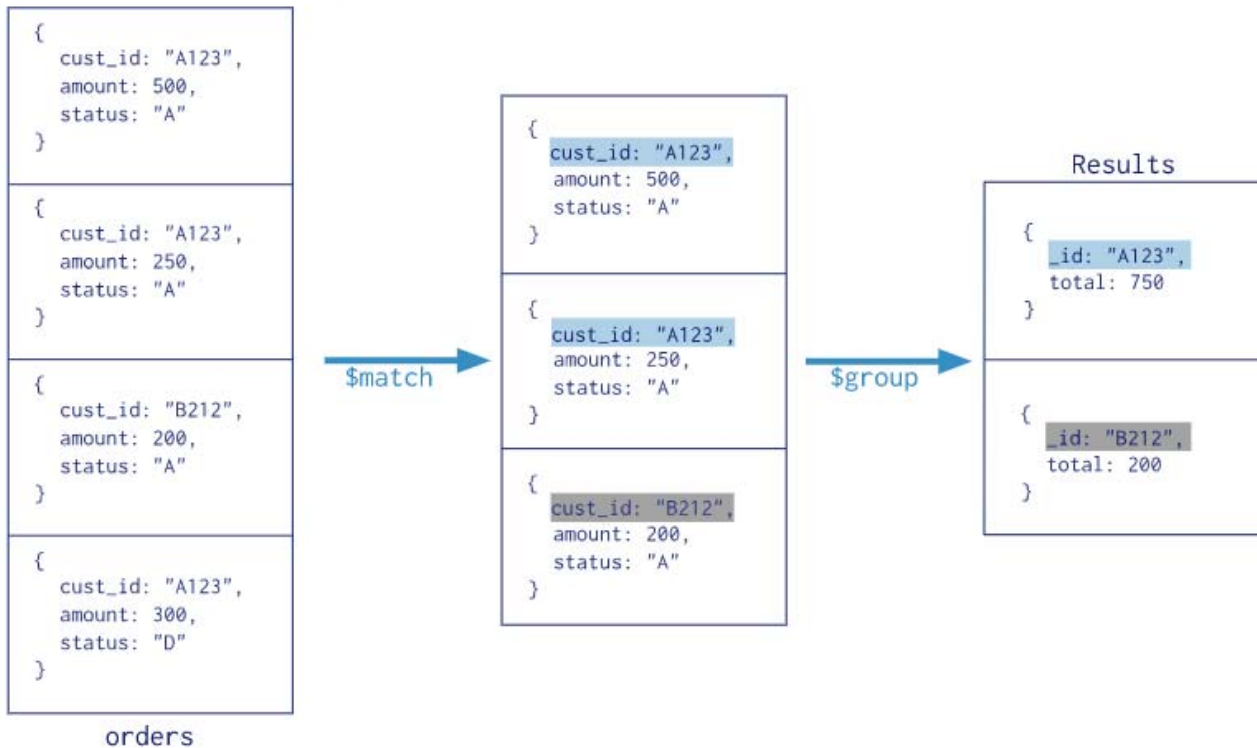
Aggregation

- For the aggregation in MongoDB, use **aggregate()** method
- Basic syntax

`db.COLLECTION_NAME.aggregate(AGGREGATE_OPERATION)`

MongoDB Aggregation Pipeline

Collection
↓
db.orders.aggregate([
 \$match stage → { \$match: { status: "A" } },
 \$group stage → { \$group: { _id: "\$cust_id", total: { \$sum: "\$amount" } } }
])



\$match stage

- Consider the **orders** collection

```
db.orders.insertMany( [  
  { cust_id: "A123", amount: 500, status: "A" },  
  { cust_id: "A123", amount: 250, status: "A" },  
  { cust_id: "B212", amount: 200, status: "A" },  
  { cust_id: "A123", amount: 300, status: "D" }  
] );
```

\$match stage

- Run the following command
`db.orders.aggregate ([
 { $match : { status : "A" } }
]);`

which is equivalent to

`db.orders.find({status:"A"})`

\$match acts as a filter to reduce the amount of documents that are given as input to the next stage

\$match stage

- More examples

```
db.orders.aggregate ( [  
    { $match : { cust_id: "A123", status : "A" } }  
] );
```

```
db.orders.aggregate ( [  
    { $match : { $or: [ {amount: {$gte: 300}} , {status : "D"}] } }  
] );
```

\$group stage

```
db.orders.aggregate ( [  
    { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }  
] );
```

Equivalent to

```
select      sum(amount) as total  
from        orders  
group by    cust_id
```

Pipeline

```
db.orders.aggregate ( [  
    { $match : { status : "A" } },  
    { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }  
] );
```

Equivalent to

```
select      cust_id, sum(amount) as total  
from        orders  
where       status= "A"  
group by    cust_id
```


Add \$sort to pipeline

```
db.orders.aggregate ( [  
    { $match : { status : "A" } },  
    { $group: { _id: "$cust_id", total: { $sum: "$amount" } } },  
    { $sort : { _id : 1 } }  
] );
```

\$count

```
db.orders.aggregate ( [  
    { $match : { status : "A" } },  
    { $group: { _id: null, count: { $sum: 1 } } }  
] );
```

The `_id` field is required for grouping and would normally contain fields from each document that we'd like to preserve. We make it null here to just count the number of documents

`$sum: 1` count the number of documents

\$avg, \$min, \$max

```
db.orders.aggregate ( [  
    { $match : { status : "A" } },  
    { $group: { _id: null, avg_amount: { $avg: "$amount" } } }  
] );
```

```
db.orders.aggregate ( [  
    { $group: { _id: "$cust_id", max_amount: { $max: "$amount" } } }  
] );
```

Pipeline Optimization

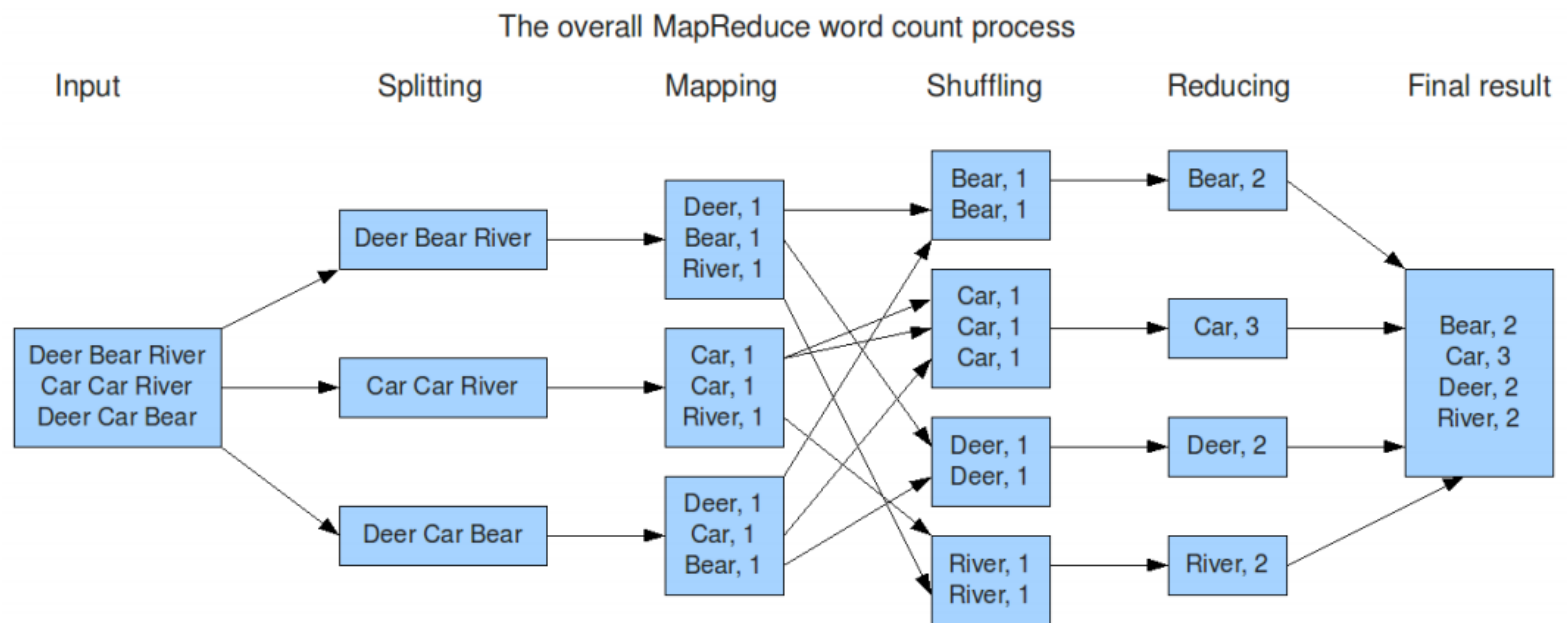
- Place the \$match as early in the aggregation pipeline as possible. Because \$match limits the total number of documents in the aggregation pipeline, earlier \$match operations minimize the amount of processing down the pipe

MapReduce

What is MapReduce

- Origin from Google, [OSDI'04]
- A simple programming model
- For large-scale data processing
- Exploits large set of commodity computers
- Executes process in distributed manner
- Two basic operations on the input
 - Map
 - Reduce

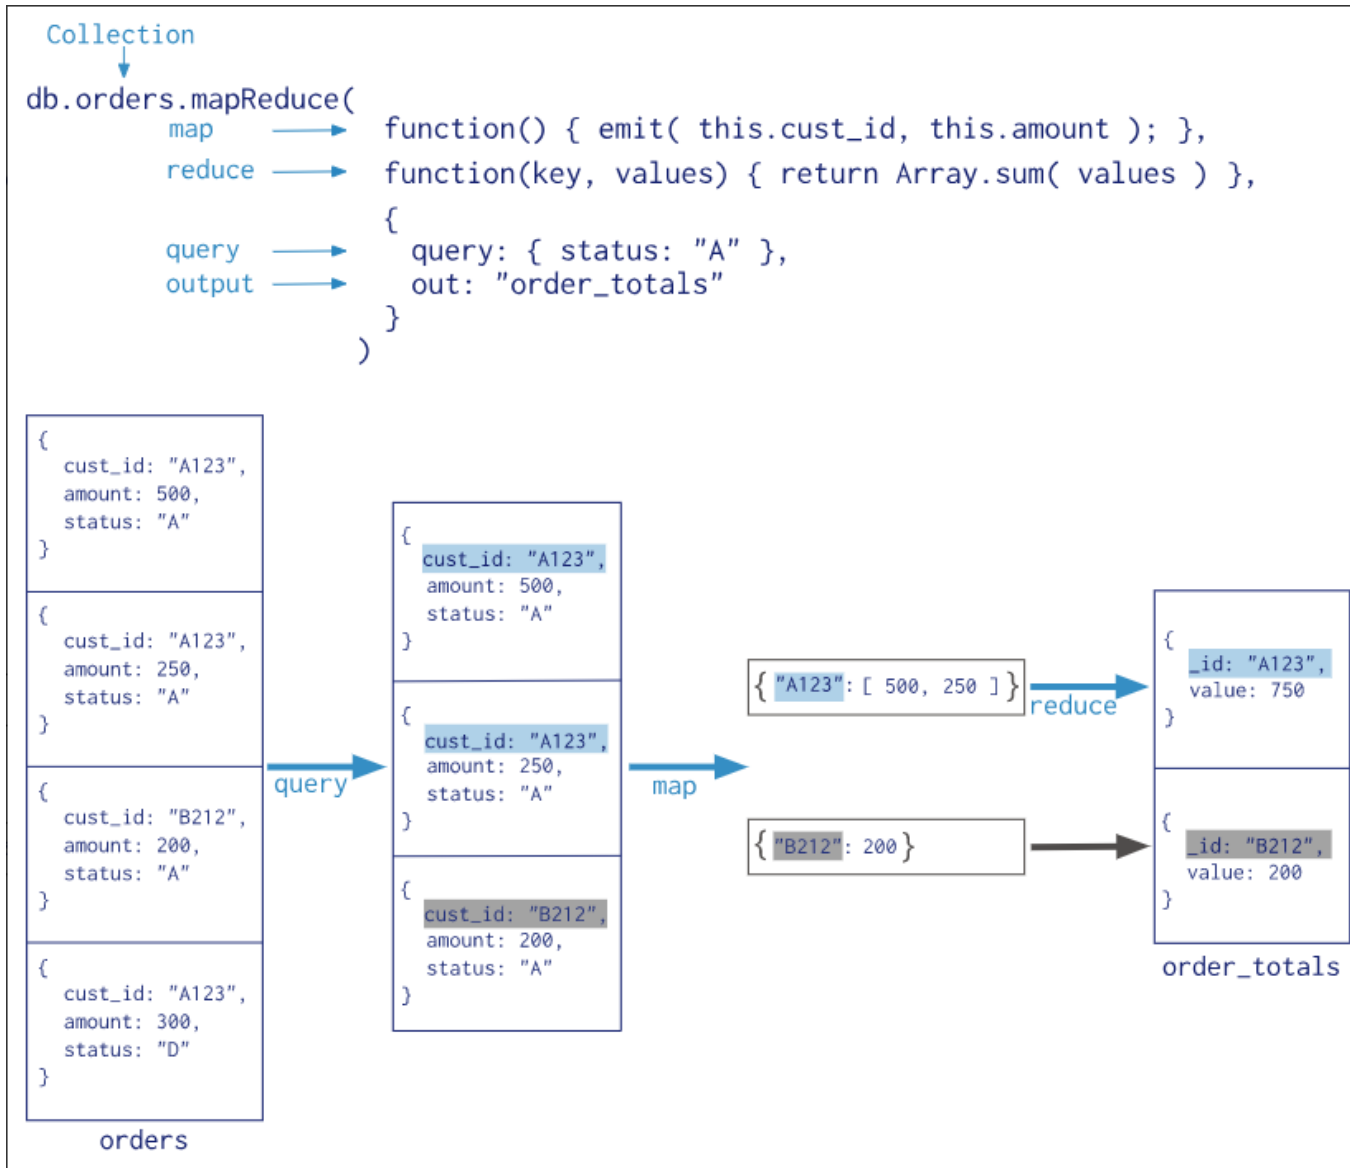
How MapReduce Works



MapReduce in MongoDB

- MongoDB also provides map-reduce operations to perform aggregation
- In the map-reduce operation
 - MongoDB applies the *map* phase to each input document (i.e. the documents in the collection that match the query condition). The map function emits key-value pairs.
 - For those keys that have multiple values, MongoDB applies the *reduce* phase, which collects and condenses the aggregated data. MongoDB then stores the results in a collection.

<https://docs.mongodb.com/manual/reference/method/db.collection.mapReduce/>



MapReduce Output

```
{
  "result" : "order_totals",           // Result stored in collection order_totals
  "timeMillis" : 659,                  // The command execution time in milliseconds
  "counts" : {
    "input" : 3,                       // The number of input documents, which is the number of
                                      // times the mapReduce command called the map function
    "emit" : 3,                       // The number of times the mapReduce command called the
                                      // emit function
    "reduce" : 1,                     // The number of times the mapReduce command called the
                                      // reduce function
    "output" : 2                      // The number of output values produced
  },
  "ok" : 1                            // A value of 1 indicates the mapReduce command ran
                                      // successfully. A value of 0 indicates an error
}
```

MapReduce in MongoDB

- MongoDB also provides map-reduce operations to perform aggregation
- In the map-reduce operation
 - MongoDB applies the *map* phase to each input document (i.e. the documents in the collection that match the query condition). The map function emits key-value pairs.
 - For those keys that have multiple values, MongoDB applies the *reduce* phase, which collects and condenses the aggregated data. MongoDB then stores the results in a collection.

MongoDB Join

\$lookup for Join

- To perform an equality match between a field from the input documents with a field from the documents of the “joined” collection, the \$lookup stage has the following syntax:

```
{  
  $lookup:  
  {  
    from: <collection to join>,  
    localField: <field from the input documents>,  
    foreignField: <field from the documents of the "from" collection>,  
    as: <output array field>  
  }  
}
```

MongoDB Data Modeling

Design Principles

- Focus on data usage
 - Data which is accessed together should stay together
 - Need to know some data access patterns / queries before designing
- Field values can be
 - Absent
 - Set to null
 - A single value
 - An array of multiple values (e.g. phone numbers of a person)
 - [6041234567, 7781234567]

Flexible Schema

- The documents in a single collection do not need to have the same set of fields and the data type for a field can differ across documents within a collection
- To change the structure of the documents in a collection, such as add new fields, remove existing fields, or change the field values to a new type, update the documents to the new structure.

Embedded Data Model

- Embedded documents capture relationships between data by storing related data in a single document structure
- Denormalized: allow applications to retrieve and manipulate related data in a single database operation
- For many use cases in MongoDB, the denormalized data model is optimal

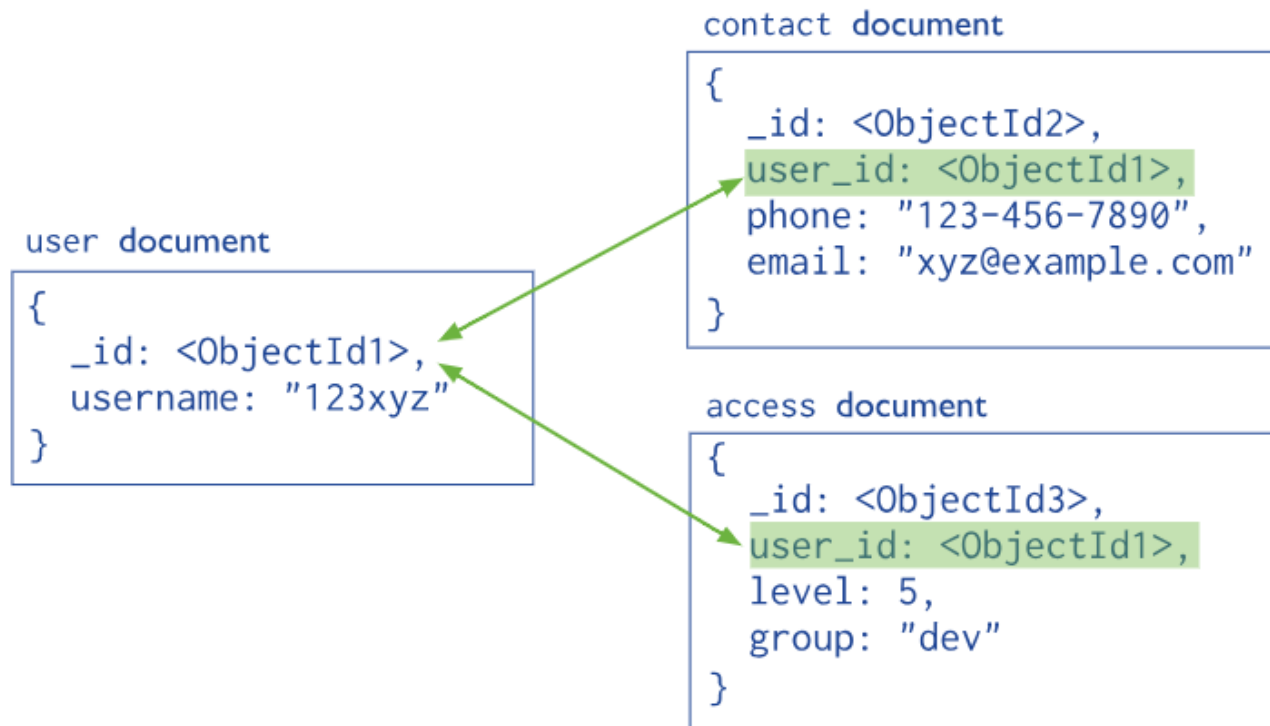
```
{
  _id: <ObjectId>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

Embedded sub-document

Embedded sub-document

Normalized Data Models (Referencing)

- Normalized data models describe relationships using references between documents.



1:1 Relationship with Embedded Documents

```
{  
  _id: "joe",  
  name: "Joe Bookreader"  
}
```

```
{  
  patron_id: "joe",  
  street: "123 Fake Street",  
  city: "Faketon",  
  state: "MA",  
  zip: "12345"  
}
```

```
{  
  _id: "joe",  
  name: "Joe Bookreader",  
  address: {  
    street: "123 Fake Street",  
    city: "Faketon",  
    state: "MA",  
    zip: "12345"  
  }  
}
```

1:M Relationship

```
{
  _id: "joe",
  name: "Joe Bookreader"
}

{
  patron_id: "joe",
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}

{
  patron_id: "joe",
  street: "1 Some Other Street",
  city: "Boston",
  state: "MA",
  zip: "12345"
}
```

```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```

M:N Relationships

- No Intersection collection (table)
- Add array in both collection
 - Need to maintain data in both for consistency

Book {

authors: [...]

}

Author {

Books: [...]

}

<http://learnmongodbthehardway.com/schema/schemabasics/>

References

https://www.tutorialspoint.com/mongodb/mongodb_aggregation.htm

<https://docs.mongodb.com/manual/>

<https://appdividend.com/2018/10/25/mongodb-aggregate-example-tutorial/>

<https://appdividend.com/2018/10/26/mongodb-mapreduce-example-tutorial/>

<http://learnmongodbthehardway.com/schema/schemabasics/>