



# Chapter 3: Introduction to SQL

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Outline

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database



# History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory in the early 1970s
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
  - SQL-86
  - SQL-89
  - SQL-92
  - SQL:1999 (language name became Y2K compliant!)
  - SQL:2003, SQL:2006, SQL:2008, SQL:2011, SQL:2016



## History (Cont.)

- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
  - Not all examples here may work on your particular system



# Data Definition Language

- The SQL data-definition language (DDL) allows the specification of information about relations, including:
  - The schema for each relation.
  - The domain of values associated with each attribute.
  - Integrity constraints
  - And as we will see later, also other information such as
    - ▶ The set of indices to be maintained for each relations.
    - ▶ Security and authorization information for each relation.
    - ▶ The physical storage structure of each relation on disk.



# Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p*,*d*)**. Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point. (ex., **numeric**(3,1), allows 44.5 to be stored exactly, but not 444.5 or 0.32)
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.
- More are covered in Chapter 4.



# BASIC SCHEMA DEFINITION



# Create Table Construct

- An SQL relation is defined using the **create table** command:

```
create table  $r$  ( $A_1 D_1, A_2 D_2, \dots, A_n D_n,$   
                (integrity-constraint1),  
                ...,  
                (integrity-constraintk))
```

- $r$  is the name of the relation
- each  $A_i$  is an attribute name in the schema of relation  $r$
- $D_i$  is the data type of values in the domain of attribute  $A_i$





# Create Table Construct (Cont.)

## ■ Example:

```
create table instructor (  
    ID           char(5),  
    name        varchar(20),  
    dept_name   varchar(20),  
    salary     numeric(8,2)  
  
);
```



# Integrity Constraints in Create Table

- **primary key**  $(A_1, \dots, A_n)$
- **foreign key**  $(A_m, \dots, A_n)$  **references**  $r$
- **not null**

*Example:*

```
create table instructor (  
    ID          char(5),  
    name        varchar(20) not null,  
    dept_name varchar(20),  
    salary      numeric(8,2),  
    primary key (ID),  
    foreign key (dept_name) references department);
```

**primary key** declaration on an attribute automatically ensures **not null**



## And a Few More Relation Definitions

```
create table student (  
    ID                varchar(5),  
    name             varchar(20) not null,  
    dept_name        varchar(20),  
    tot_cred         numeric(3,0),  
    primary key (ID),  
    foreign key (dept_name) references  
    department);
```



## And a Few More Relation Definitions (Cont.)

```
create table takes (  
    ID                varchar(5),  
    course_id        varchar(8),  
    sec_id            varchar(8),  
    semester         varchar(6),  
    year              numeric(4,0),  
    grade             varchar(2),  
    primary key (ID, course_id, sec_id, semester,  
year) ,  
    foreign key (ID) references student,  
    foreign key (course_id, sec_id, semester,  
year) references section);
```



## And more still

```
create table course (  
    course_id      varchar(8),  
    title          varchar(50),  
    dept_name     varchar(20),  
    credits        numeric(2,0),  
    primary key (course_id),  
    foreign key (dept_name) references  
    department);
```



# Updates to tables

## ■ Insert

- **insert into** *instructor* **values** ('10211', 'Smith', 'Biology', 66000);

## ■ Delete

- Remove all tuples from the *student* relation
  - ▶ **delete from** *student*

## ■ Drop Table

- **drop table** *r*

## ■ What is the difference between delete and drop?



# Updates to tables (Cont.)

## ■ Alter

### ● **alter table $r$ add $A$ $D$**

- ▶ where  $A$  is the name of the attribute to be added to relation  $r$  and  $D$  is the domain of  $A$ .
- ▶ All exiting tuples in the relation are assigned *null* as the value for the new attribute.

### ● **alter table $r$ drop $A$**

- ▶ where  $A$  is the name of an attribute of relation  $r$
- ▶ Dropping of attributes not supported by many databases. Only allow to drop the whole table.
- ▶ Can also drop constraint



# Basic Query Structure

- A typical SQL query has the form:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

- $A_i$  represents an attribute
  - $R_i$  represents a relation
  - $P$  is a predicate.
- The result of an SQL query is a relation.
    - But, may have duplicates, not a relation according to formal mathematical definition





# The select Clause

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:  
**select** *name*  
**from** *instructor*



## The select Clause (Cont.)

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after select.
- Find the department names of all instructors, and remove duplicates

```
select distinct dept_name  
from instructor
```

- The keyword **all** specifies that duplicates should not be removed.

```
select all dept_name  
from instructor
```



# The select Clause (Cont.)

- An asterisk in the select clause denotes “all attributes”

**select \***  
**from** *instructor*

<i>name</i>
Srinivasan
Wu
Mozart
Einstein
El Said
Gold
Katz
Califieri
Singh
Crick
Brandt
Kim

**Figure 3.2** Result of “select *name* from *instructor*”.

<i>dept_name</i>
Comp. Sci.
Finance
Music
Physics
History
Physics
Comp. Sci.
History
Finance
Biology
Comp. Sci.
Elec. Eng.

**Figure 3.3** Result of “select *dept\_name* from *instructor*”.



## The select Clause (Cont.)

- The **select** clause can contain arithmetic expressions involving the operation, +, −, \*, and /, and operating on constants or attributes of tuples.

- The query:

```
select ID, name, salary/12  
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

- Can rename “*salary/12*” using the **as** clause:

```
select ID, name, salary/12 as monthly_salary
```



# The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```



## The where Clause (Cont.)

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
  - To find all instructors in Comp. Sci. dept with salary > 80000

```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary >
80000
```

- Comparisons can be applied to results of arithmetic expressions.



# The from Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.

- Find the Cartesian product *instructor X teaches*

**select \***  
**from** *instructor, teaches*

- generates every possible instructor – teaches pair, with all attributes from both relations.
  - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra).



# Cartesian Product

*instructor*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

*teaches*

<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

<i>Inst.ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>teaches.ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Pinance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Pinance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Pinance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...





# Examples

- Find the names of all instructors who have taught some course and the course\_id

```
select name, course_id  
from instructor, teaches  
where instructor.ID = teaches.ID
```

- Find the names of all instructors in the Art department who have taught some course and the course\_id

```
select name, course_id  
from instructor , teaches  
where instructor.ID = teaches.ID and instructor.  
dept_name = 'Art'
```



# Understanding SQL Query

- In general, the meaning of an SQL query can be understood as follows:
  - Generate a Cartesian product of the relations listed in the **from** clause
  - Apply the predicates specified in the **where** clause on the result of Step 1
  - For each tuple in the result of Step 2, output the attributes (or results of expressions) specified in the **select** clause

Note: The above sequence of steps helps make clear what the result of an SQL query should be, **not** how it should be executed. A real implementation of SQL optimize evaluation by generating (as far as possible) only elements of the Cartesian product that satisfy the **where** clause predicates.



# The Rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

*old-name as new-name*

- Example:

```
select name as instructor_name, course_id  
from instructor, teaches  
where instructor.ID = teaches.ID;
```

```
select T.name, S.course_id  
from instructor as T, teaches as S  
where T.ID = S.ID;
```



## The Rename Operation (Cont.)

- Find the names of all instructors who have a higher salary than at least one instructor in 'Comp. Sci'.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept_name = 'Comp.  
Sci.'
```

- *T* and *S* referred to as correlation name, table alias, correlation variable, or tuple variable
- Keyword **as** is optional and may be omitted  
*instructor* **as** *T*  $\equiv$  *instructor T*



# Join

- Matches attributes with the same value
- Consider the query “For all instructors in the university who have taught some course, find their names and the course ID of all courses they taught”, we can write as

**select** *name, course id*

**from** *instructor, teaches*

**where** *instructor.ID = teaches.ID;*



## Join (Cont.)

- It is equivalent to  
**select** *name, course id*  
**from** *instructor* **join** *teaches*  
**on** *instructor.ID = teaches.ID;*



# String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring.
  - underscore ( \_ ). The \_ character matches any character.



## String Operations (Cont.)

- Find the names of all instructors whose name includes the substring “dar”.

```
select name  
from instructor  
where name like '%dar%'
```

- Match the string “100%”

```
like '100 \%' escape '\'
```

in that above we use backslash (\) as the escape character.





# String Operations (Cont.)

- Patterns are case sensitive.
- Pattern matching examples:
  - 'Intro%' matches any string beginning with "Intro".
  - '%Comp%' matches any string containing "Comp" as a substring.
  - '\_\_\_' matches any string of exactly three characters.
  - '\_\_\_ %' matches any string of at least three characters.



# String Operations (Cont.)

- SQL supports a variety of string operations such as
  - Concatenation (CONCAT)
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.

<https://docs.microsoft.com/en-us/sql/t-sql/functions/string-functions-transact-sql?view=sql-server-2017>



# Ordering the Display of Tuples

- List in alphabetic order the names of all instructors

```
select distinct name  
from instructor  
order by name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

- Example: **order by** *name* **desc**

- Can sort on multiple attributes

- Example: **order by** *dept\_name*, *name*

- **order by** *salary* **desc**, *name* **asc**;



# Where Clause Predicates

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq$  \$90,000 and  $\leq$  \$100,000)

```
select name  
from instructor  
where salary between 90000 and 100000
```



## ■ Consider the following problems

- Find courses (just course id information) that ran in Fall 2009 or in Spring 2010
- Find courses that ran in both Fall 2009 and Spring 2010
- Find courses that ran in Fall 2009 but not in Spring 2010



# Set Operations

- Find courses that ran in Fall 2009 or in Spring 2010

**(select *course\_id* from *section* where *semester* = 'Fall' and *year* = 2009)**  
**union**

**(select *course\_id* from *section* where *semester* = 'Spring' and *year* = 2010)**

- Find courses that ran in both Fall 2009 and Spring 2010

**(select *course\_id* from *section* where *semester* = 'Fall' and *year* = 2009)**  
**intersect**

**(select *course\_id* from *section* where *semester* = 'Spring' and *year* = 2010)**

- Find courses that ran in Fall 2009 but not in Spring 2010

**(select *course\_id* from *section* where *semester* = 'Fall' and *year* = 2009)**  
**except**

**(select *course\_id* from *section* where *semester* = 'Spring' and *year* = 2010)**



## Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary.

**(select distinct *salary***  
**from *instructor*)**

**except**

**(select max(salary)**  
**from *instructor*)**



## Set Operations (Cont.)

- Set operations **union**, **intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**
- Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:
  - $m + n$  times in  $r$  **union all**  $s$
  - $\min(m, n)$  times in  $r$  **intersect all**  $s$
  - $\max(0, m - n)$  times in  $r$  **except all**  $s$





# Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```



# Null Values and Three Valued Logic

- Three values – *true*, *false*, *unknown*
- Any comparison with *null* returns *unknown*
  - Example:  $5 < null$  or  $null <> null$  or  $null = null$
- Three-valued logic using the value *unknown*:
  - OR:  $(unknown \text{ or } true) = true$ ,  
 $(unknown \text{ or } false) = unknown$   
 $(unknown \text{ or } unknown) = unknown$
  - AND:  $(true \text{ and } unknown) = unknown$ ,  
 $(false \text{ and } unknown) = false$ ,  
 $(unknown \text{ and } unknown) = unknown$
  - NOT:  $(\text{not } unknown) = unknown$
  - “*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*



# Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values



## Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department

```
select avg (salary)  
from instructor  
where dept_name= 'Comp. Sci.';
```

- Find the total number of instructors who teach a course in the Spring 2010 semester

```
select count (distinct ID)  
from teaches  
where semester = 'Spring' and year = 2010;
```

- Find the number of tuples in the *course* relation

```
select count (*)  
from course;
```



# Aggregate Functions – Group By

- Find the average salary of instructors in each department
  - **select** *dept\_name*, **avg** (*salary*) **as** *avg\_salary*  
**from** *instructor*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

<i>dept_name</i>	<i>avg_salary</i>
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000



## Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list
  - */\* erroneous query \*/*  
**select** *dept\_name*, *ID*, **avg** (*salary*)  
**from** *instructor*  
**group by** *dept\_name*;



# Aggregate Functions – Having Clause

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups



# Null Values and Aggregates

## ■ Total all salaries

```
select sum (salary )  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount

## ■ All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes

## ■ What if collection has only null values?

- count returns 0
- all other aggregates return null





# Exercise

- Using the files provided to generate the University schema
- Write the following queries
  - Find the titles of courses in the Comp. Sci. department that have 3 credits.
  - Find the highest salary of any instructor.
  - Find all instructors earning the highest salary
  - Find the enrollment of each section that was offered in Spring 2009



# Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.



# Subqueries in the Where Clause



# Subqueries in the Where Clause

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality



# Set Membership

- Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id  
from section  
where semester = 'Fall' and year = 2009 and  
       course_id in (select course_id  
                        from section  
                        where semester = 'Spring'  
                        and year = 2010);
```



## Set Membership (Cont.)

- Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
       course_id not in (select course_id
                           from section
                           where semester = 'Spring' and
                           year= 2010);
```



## Set Comparison – “some” Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Physics department.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept_name = Physics';
```

Same query using > **some** clause

```
select name  
from instructor  
where salary > some (select salary  
                        from instructor  
                        where dept_name = 'Physics');
```



# Definition of “some” Clause

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$  (read:  $5 < \text{some tuple in the relation}$ )

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However,  $(\neq \text{some}) \equiv \text{not in}$





## Set Comparison – “all” Clause

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Computer Science department.

```
select name  
from instructor  
where salary > all (select salary  
                        from instructor  
                        where dept_name = 'Comp. Sci.');
```



# Definition of “all” Clause

(5 < **all**

0
5
6

) = false

(5 < **all**

6
10

) = true

(5 = **all**

4
5

) = false

(5 ≠ **all**

4
6

) = true (since 5 ≠ 4 and 5 ≠ 6)

(≠ **all**) ≡ **not in**

However, (= **all**) ≠ **in**



# Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$



# Use of “exists” Clause

- Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
      exists (select *
              from section as T
              where semester = 'Spring' and year = 2010
              and S.course_id = T.course_id);
```

- **Correlation name** – variable *S* in the outer query
- **Correlated subquery** – the inner query



# Use of “exists” Clause

- Find all courses taught in the Fall 2009 semester but not in the Spring 2010 semester.



# Subqueries in the From Clause



# Subqueries in the From Clause

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name) as dept_average
where avg_salary > 42000;
```

- Note that we do not need to use the **having** clause



## With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the budget over the average  
**with** *avg\_budget(value)* **as**  
    (**select avg**(*budget*)  
    **from** *department*)  
**select** *department.dept\_name*  
**from** *department, avg\_budget*  
**where** *department.budget* > *avg\_budget.value*;





# Subqueries in the Select Clause



# Subquery in SELECT

- List all departments along with the number of instructors in each department

```
select dept_name,  
      (select count(*)  
      from instructor  
      where department.dept_name = instructor.dept_name)  
      as num_instructors  
from department;
```